

# Bayesian Simulation

Aleksei Sorokin, asorokin@hawk.iit.edu, A20394300

## Load Data

```
df <- read.csv('data/all_coaches.csv')
data_dict <- list(
  N='Name',
  GR='Games Relative',
  WP='Win Loss Percentage',
  PGR='Playoff Games Relative',
  PWP='Playoff Win Percentage',
  CC='Conference Championships',
  C='Championships',
  HOF='Hall of Fame',
  S='Sport')
names(df) <- names(data_dict)
head(df)
```

##		N	GR	WP	PGR	PWP	CC	C	HOF	S
## 1	AJ Hinch	6.308642	0.5580000	4.5454545	0.560	2	1	0	baseball	
## 2	Aaron Boone	2.000000	0.6270000	1.2727273	0.500	0	0	0	baseball	
## 3	Aaron Kromer	0.375000	0.3330000	0.0000000	0.000	0	0	0	football	
## 4	Abe Gibrón	2.625000	0.2740000	0.0000000	0.000	0	0	0	football	
## 5	Adam Gase	4.000000	0.4690000	0.3333333	0.000	0	0	0	football	
## 6	Adam Oates	1.585366	0.5752212	0.4375000	0.429	0	0	0	hockey	

## Split into train test datasets

```
set.seed(7)
train_frac <- 4/5
df_HOF_1 <- df[df$HOF==1,]
df_HOF_0 <- df[df$HOF==0,]
HOF_1_train <- sample(1:nrow(df_HOF_1), floor(nrow(df_HOF_1)*train_frac))
HOF_0_train <- sample(1:nrow(df_HOF_0), floor(nrow(df_HOF_0)*train_frac))
df_train <- rbind(df_HOF_1[HOF_1_train,], df_HOF_0[HOF_0_train,])
df_test <- rbind(df_HOF_1[-HOF_1_train,], df_HOF_0[-HOF_0_train,])
l1 <- sprintf('Train Fraction: %.2f', train_frac)
l2 <- sprintf('Hall of Fame Coaches: %d. (Train %d , Test %d)',
  nrow(df_HOF_1), length(HOF_1_train), nrow(df_HOF_1)-length(HOF_1_train))
l3 <- sprintf('Non Hall of Fame Coaches: %d. (Train %d , Test %d)',
  nrow(df_HOF_0), length(HOF_0_train), nrow(df_HOF_0)-length(HOF_0_train))
l4 <- sprintf('Overall: (Train %d , Test %d)', nrow(df_train), nrow(df_test))
cat(sprintf('%s\n%s\n%s\n%s\n', l1, l2, l3, l4))
```

```
## Train Fraction: 0.80
## Hall of Fame Coaches: 256. (Train 204 , Test 52)
## Non Hall of Fame Coaches: 1664. (Train 1331 , Test 333)
## Overall: (Train 1535 , Test 385)
```

## Logistic Regression with only GR

using lm from R

```
model_1 <- glm(HOF ~ GR, data=df_train, family="binomial")
summary(model_1)
```

```
##
## Call:
## glm(formula = HOF ~ GR, family = "binomial", data = df_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3277  -0.4978  -0.4364  -0.4141   2.2491
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.44816    0.10628 -23.035  <2e-16 ***
## GR           0.11180    0.01188   9.414  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1203.0  on 1534  degrees of freedom
## Residual deviance: 1115.3  on 1533  degrees of freedom
## AIC: 1119.3
##
## Number of Fisher Scoring iterations: 4
```

Posterior Simulation from non-informative prior

```
# constants
n_sims <- 1000
x <- df_train[['GR']]
y <- df_train[['HOF']]
length(x) == length(y)

## [1] TRUE


n <- length(y)
n_mesh <- 100
# compute posterior density on grid
beta0 <- seq(-3,-2,length=n_mesh)
beta1 <- seq(0.05,.2,length=n_mesh)
logl <- function(b0,b1){
  z <- b0+b1*x
  theta <- exp(z)/(1+exp(z))
  d <- sum( y*z+log(1-theta) )
  return(d)}
log_post_dens <- matrix(NA, n_mesh, n_mesh)
```

```

for (i in 1:n_mesh){
  for (j in 1:n_mesh){
    log_post_dens[i,j] <- log1(beta0[i],beta1[j])}
post_dens <- exp(log_post_dens - max(log_post_dens))
# contour plot
contour(beta0, beta1, post_dens, levels=c(.001,.01,.1,.3,.5,.7,.9),
        xlab='beta0', ylab='beta1', main='Posterior Density Contour')

```

**Posterior Density Contour**



```

# simulations
beta0_density <- rowSums(post_dens)
beta0_idx <- sample(1:n_mesh, n_sims, replace=T, prob=beta0_density)
b0_sims <- beta0[beta0_idx]
b1_sims <- rep(NA,n_sims)
for (i in 1:n_sims){
  beta1_density_i <- post_dens[beta0_idx[i],]
  b1_sims[i] <- exp(sample(beta1, 1, prob=beta1_density_i))}
# intervals
I_b0 <- quantile(b0_sims,c(0.05,0.95))
I_b1 <- quantile(b1_sims,c(0.05,0.95))
# outputs
cat(sprintf('90%% interval estimate for a: (%.2f,%.2f)\n',I_b0[1],I_b0[2]))

## 90% interval estimate for a: (-2.64,-2.28)
cat(sprintf('90%% interval estimate for b: (%.2f,%.2f)\n',I_b1[1],I_b1[2]))

## 90% interval estimate for b: (1.10,1.14)

```